Do You Know About My Nation? Investigating Multilingual Language Models' Cultural Literacy Through Factual Knowledge

Anonymous ACL submission

Abstract

Most multilingual question-answering benchmarks, while covering a diverse pool of languages, do not factor in regional diversity in the information they capture and tend to be Western-centric. This introduces a significant gap in fairly evaluating multilingual models' comprehension of factual information from diverse geographical locations. To address this, we introduce XNationQA to investigate the *cultural literacy* of multilingual LLMs. XNationQA encompasses questions on the geography, culture, and history of nine countries, containing a total of 49,280 questions in seven languages. We benchmark eight standard multilingual LLMs on XNationQA. Our analyses uncover a considerable discrepancy in the model's accessibility to culturally specific facts across languages. Notably, we often find that a model demonstrates greater knowledge of cultural information in English than in the dominant language of the respective culture. The models show better performance in Western languages, though this does not translate to their being more literate for Western countries, which is counter-intuitive. Furthermore, we observe that models have a very limited ability to transfer knowledge across languages, particularly evident in open-source models.

1 Introduction

Multilingual Large Language Models (LLMs) (Üstün et al., 2024; Achiam et al., 2023) show impressive performance on many lanugages across varied tasks. However, the best practices to evaluate them remain contested (Ahuja et al., 2023; Hada et al., 2023; Saha et al., 2023), with many criticising Western-centric (often Anglosphere) evaluations (Held et al., 2023). For example, Faisal et al. (2021) demonstrated how multiple purportedly multilingual question answering (QA) datasets disproportionately cover USA-related concepts and entities, with a similar question distribution to mono-

lingual English benchmarks. This is because multilingual QA benchmarks are often derived from western-centric English (Longpre et al., 2021; Kassner et al., 2021; Dumitrescu et al., 2021). They, thus, fail to consider the cultural contexts where these languages are spoken (Naous et al., 2023). Blevins et al. (2022) did propose guidance techniques to create more inclusive benchmarks, to the best of our knowledge, there is no multilingual large-scale QA dataset which explicitly balances its distribution across a set of diverse cultures. LLMs are clearly able to access factual knowledge in multiple languages (Jiang et al., 2020; Kassner et al., 2021), but the relationship between language and information locality is poorly understood. To this end, we produce a parallel knowledge test, *consisting of regionally differing factual information*, balanced across a set of related test cultures and languages, motivated from the education research literature.

In education theory, *cultural literacy* refers to the shared corpus of translinguistic cultural knowledge within a community, polity, or society on which efficient communication is built (Hirsch, 1983). By *translinguistic*, Hirsch (1983) meant that the information in the corpus is language-agnostic in its meaning. An interesting study conducted by Steffensen et al. (1979) in this regard found that American English speakers performed worse than Indian English speakers on a reading comprehension test of a story describing an Indian wedding, and vice versa on a story of an American one at equivalent reading levels because the context was more familiar to the speakers of the corresponding culture. Multilingual LLMs strive to create technology that is inclusive of a diverse set of linguistic group (Blasi et al., 2021). Therefore, they need to be culturally literate if they are to consistently answer translinguistic regional questions about specific topics consistently across languages and nations.

Figure 1: The working example aims to evaluate the cultural literacy of an LLM about Spain across English and Spanish languages. While the questions related to *war* and *monument* are answered correctly in both languages, the model fails to do so for the question related to *political leader* and could only solve the *national park* question in Spanish, indicating a lack of cultural literacy about these topics and across languages.

Using cultural literacy as a framing device, we investigate multilingual, multicultural knowledge in multilingual language models by answering these research questions:

RQ1: *What disparities exist in the cultural literacy of various LLMs about different countries?*

RQ2: *How transferable is the cultural literacy of LLMs across different languages?*

RQ3: *How consistent is the LLM in recalling factoids of nations across different languages?*

RQ4: *Does a systematic relationship between country and language exist in LLM cultural literacy?*

To this end, we create XNationQA, a test set designed to assess cultural literacy through QA pairs that evaluate curricular knowledge for nine countries (Japan, India, China, Germany, Spain, Russia, Mexico, the USA, the UK), fully translated into seven languages (Russian, Hindi, German, Spanish, Japanese, English, Chinese). Given culture's nuanced and context-dependent nature, which can vary across communities, languages, and sociopolitical environments, it is inherently challenging to define and quantify. To operationalise cultural literacy in our study, we draw on prior work in the monolingual (English-centric) setting that measures geographic and cultural erosion in language models via factual recall (Schwöbel et al., 2023; Zhou et al., 2022). Following this approach, we

assess a model's cultural literacy about a nation through its performance on questions concerning historically and geographically grounded facts specific to that nation. We categorize these facts into four domains:(1) *wars*, or the year of armed conflicts undertaken by the country, (2) *national parks*, or the site of nationally designated protected wilderness areas in each country, (3) *leaders*, or the date of birth of heads of state of the country, and (4) *monuments*, or UNESCO world heritage site locations in the country.This framing allows us to systematically evaluate a model's knowledge of culturally embedded facts, serving as a proxy for its cultural literacy. Figure 1 shows a working example of our task. We prompt the model to answer questions related to monuments, war, leaders, and national parks. It is evident that the model has a disparity in answering the question across languages, as observed in the national park and monuments question. In the leader-related question, the model fails to answer the facts.

We extensively evaluate nine commonly-used multilingual LLMs on XNationQA and find that — (i) There is a disparity in the cultural literacy of LLMs across different languages, with models exhibiting varying knowledge. Models tend to perform the best in English, followed by other Western languages (German, Spanish, and Russian). (ii)

2

Most LLMs show poor transference of facts across languages, with open-source models struggling to answer faithfully in all languages. When considering only Western languages, models tend to be more faithful, indicating a knowledge coverage disparity between Western and non-Western languages. (iii) LLMs exhibit varying performance across nations, with different models exhibiting different levels of literacy about different nations. Surprisingly, a culture's native language is not the best-performing language in the domain, with Western languages performing the best. (iv) We also find that while Western languages are preferred by the model, they don't show higher performance in Western countries over non-western ones, in fact, nations that seem to be more literate about India, China, and Japan than Germany, Russia or Spain.[1]

2   Related Work and Motivation

Multilingual LLMs.   Since the advent of pre-trained transformer-based language models (Devlin et al., 2018), there has been a constant effort to develop multilingual LLMs that can understand and reason in multiple languages. These variants are trained on unsupervised training objectives using large multilingual corpora such as Oscar (Abadji et al., 2022), mC4 (Xue et al., 2021), and CulturaX (Nguyen et al., 2023), which are not parallel across languages. Therefore, the multilingual generalization ability of these models is a by-product of their ability to project different languages into a common representation space (Artetxe et al., 2019; Blevins et al., 2022). This representation ability directly depends on the datasets they are trained on (Deshpande et al., 2022). However, studies have shown that they are biased toward Western concerts because of their training mix (Naous et al., 2023; Cao et al., 2023).

Multilingual benchmarks.   Multiple multilingual benchmarks spanning various NLP tasks (Dac Lai et al., 2023) have been created. However, these benchmarks are commonly derived from monolingual English benchmarks (Clark et al., 2020; Dumitrescu et al., 2021) and, hence, tend to be biased in their coverage. Our initial analysis using multilingual QA datasets reveals these problems. For example, MKQA (Longpre et al., 2021), a popular closed-book QA dataset, is generated by translating $10,000$ samples from

Google's Natural Questions dataset (Kwiatkowski et al., 2019) without any cultural or regional considerations. We find the resultant dataset to contain questions about popular Hollywood shows like 'Modern Family' and 'Rick and Morty', while it does not cover Bollywood or other regional shows. Hence, these datasets fail to highlight if the model is culturally literate across languages and regions. In contrast XNationQA is large ($49,280$ questions) and covers multiple specific topics in parallel sets for nine nations.

Factual and cultural knowledge in multilingual models.   Several recent studies have begun to probe cultural knowledge in multilingual settings(Fung et al., 2024; Shi et al., 2024). Keleg and Magdy (2023) explores multilingual factual knowledge by mining Wikidata triples in four languages, but its reliance on language-linked Wikidata labels restricts coverage to entities available in all languages—overlooking culturally important items like Mahatma Gandhi Marine National Park, hence lacking uniform representation. XNationQA addresses such limitations by employing translation toolkits to construct a parallel corpus across seven languages, ensuring equitable coverage of historically and culturally significant content.   BLEnD (Myung et al., 2024) focuses on folk and everyday-life sourced from native speakers across 16 countries. CaLM-QA (Arora et al., 2024) studies culturally grounded long-form questions, testing whether LLMs can generate detailed answers rooted in cultural contexts. While both BLEnD and CaLM-QA emphasise cultural specificity, they differ from our work in scope and evaluation paradigm: BLEnD focuses on subjective, daily-life knowledge, and CaLM-QA necessitates human evaluation, and a high budget for such evaluation, due to its generative format. In contrast, XNationQA targets factual and historically anchored knowledge that is widely documented (e.g., wars, national parks, leaders), using a multiple-choice setup that enables scalable and easy benchmarking and evaluation for future studies.

3   XNationQA Dataset

XNationQA is a parallel multilingual dataset encompassing factual knowledge related to nine countries and seven languages. XNationQA is designed to evaluate the cultural literacy of LLMs across languages by testing their knowledge of nation-specific information. Each instance within

XNationQA includes an objective-type question focusing on a specific domain, with one correct and three incorrect options. In total, the dataset contains 49, 280 questions, spanning 1, 760 factual entities. Figure 3 shows the distribution of the entities across nine countries we covered.

Mining entities. We began by first mining factual entities specific to a country. We used English Wikipedia pages to obtain lists for four specific domains: (i) leaders of a country, (ii) national parks in a country, (iii) UNESCO sites in the nation, and (iv) wars the nation has participated in. This method was applied to extract entities for nine countries: China, Germany, India, Japan, Mexico, Russia, Spain, the UK, and the USA (see Appendix A.3 for details.)

After extracting the entity list, we used Wikidata to mine additional information. This included information about the start year of a war, the birth year of a leader, and the administrative location of a national park and the UNESCO site. Information about the entities without an appropriate Wikidata entry was filled out manually. This process produced the final entity-answer pairs used to construct the questions.

Constructing question templates and options. After extracting a nation's entity-answer pair, we manually created four prompt templates for each domain and country in English. Each domain's template was designed to ask a specific question, either about the year it began or its location.

To generate the desired objective-type questions, we generated three incorrect answers by sampling three random administrative areas of a country where the UNESCO site and national park are not located. Additionally, for year-type questions, we generated incorrect answers by randomly adding or subtracting a number between 5 and 10 from the correct year. These incorrect answers, along with the correct answer, are paired with the manually created question templates associated with an entity to form multiple-choice questions. Hence for each entity, we get four prompts. This approach is designed to account for any prompt sensitivity or bias in the LLM.

Expansion to other languages. We use translation toolkits to create the parallel multilingual corpus. This is done because not all entities in wiki data have labels in multiple languages. Hence to create XNationQA all the generated templates and entity-option pairs were translated using Google Translate and GPT-4, respectively, into Hindi, Spanish, Man-

Score	DE	ES	HI	RU	JA	ZH	AVG
Cos-sim	0.94	0.94	0.87	0.90	0.87	0.87	0.89
BLEU	50.80	67.45	68.48	43.80	52.30	45.64	54.74
H-Eval	4.85	4.75	4.57	4.25	4.50	4.41	4.55

Table 1: Validation of dataset quality using Human evaluation (H-Eval), BLEU score, and Cosine-similarity (Cos-sim). The queries seem to be semantically aligned. (see Table 4 for ISO codes)

darin, Japanese, Russian, and German. We filled out the entity and the options in the templates in their respective languages to generate the relevant objective question. This resulted in a geographically diverse, parallel multilingual factual knowledge corpus, XNationQA (c.f. Table 5).

Validation of translation. To validate the quality of XNationQA we employ back-translation (Miyabe and Yoshino, 2015) and semantic similarity to ensure the quality of the proposed dataset (c.f. Table 1).

1. **Semantic Similarity:** To ensure that queries in different languages have the same semantic meaning as their English counterparts, we compute the cosine similarity between the English queries and their translations. We use multilingual sentence transformers to extract embeddings, which are then used to compute cosine similarity. Our dataset has an average similarity score of 0.89, suggesting cross-lingual consistency and meaning preservation.

2. **Back Translation:** In this study, we randomly selected 1000 queries from XNationQA for each of the translated languages, 6000 queries in total, across the various topics in our dataset. These queries were back-translated into English using Google Translate and then compared to their original English versions. We observed an average BLEU score of 54.74 which indicates that the translations preserve the concepts of the original queries.

3. **Human Eval:** To assert the grammatical correctness and overall clarity of the dataset, 1000 queries(same as in Back Translation) from each language in XNationQA, 6000 queries in total, were evaluated by language experts. The experts assigned a score of 1 to 5 based on grammar, fluency, and coherence to assess the quality of the dataset. We found that all samples at least got a score of 4 with an average score of 4.55, indicating the

high quality of our dataset.(see Appendix B for details.)

4   Problem Definition and Experiment

Our dataset D spans L languages and covers a set of nation-specific factual entities E. For each entity $e \in E$ in our dataset, we have a set of four manually created questions q^l and options o^l where language $l \in L$; the task is to generate the correct answer a^l from the options. To do so, we formalize our prompting setup as generating the output $\hat{y}^l$ conditioned upon the question and option $\hat{y}^l = \arg\max P(y|q^l \oplus o^l)$. The goal is to match the correct answer a^l across all languages l and entities e, to have transferable cultural literacy

We experiment with eight commonly-used instruction-tuned multilingual language models, specifically 7 and 13-billion versions of LLaMA-2 (Touvron et al., 2023), 8-billion Meta-LLaMA-3-Instruct (Dubey et al., 2024), 7-billion Bloomz (Yong et al., 2023), 7-billion Mistral (Jiang et al., 2023), 7-billion Mixtral (Jiang et al., 2024), 13-billion Aya (Üstün et al., 2024) and GPT-4 (Achiam et al., 2023).For the GPT-4 model, due to budgetary constraints, we sample one question for each entity in every language. In total we evaluated GPT-4 on 12320 questions. These models have different mixtures of languages in their training corpus Refer to Appendix A.4 for more information.

5   Results and Analysis

Models' cultural literacy across languages. Table 2 presents the accuracy of LLMs across languages, averaged across countries. GPT-4 demonstrates the highest performance with an average accuracy of 72%, followed by Mixtral-8X7B with 60%, LLaMA-3-8B-Instruct with 59%, Mistral-7B-Instruct with 42.6%, Llama-2-7B-chat with 41%, Llama-2-13B-chat with 40%, Aya with 32% and Bloomz with 30%. Models seem to perform better at answering questions about *monuments* and *national parks* than questions on *wars* or *leaders*. Appendix A.7 for deeper model-wise analysis.

As evident from Table 2, LLMs have a varying degree of cultural literacy across languages, leading to noteworthy variation in performance with language change. English tends to be the best-performing language in cultural literacy. Surprisingly, Bloomz and Aya models, which have been trained specifically for multilingual alignment, supporting 46 and 101 languages respectively, under-

perform compared to models from LLaMA and Mistral families of comparable size, which have a much smaller span of languages in their pre-training data. The former models also show near-random performance on questions from *wars* and *leaders* domains, where recalling specific years is required, whereas they recall locations better. To further study the disparity in a model's performance between languages on facts about the same country, we analyze the standard deviation of the accuracy of the models in different languages. Refer to the Appendix C for its analysis.

Western vs non-Western. We also note a clear divide in performance between Western languages (English, German, Spanish, Russian) and non-Western languages (Hindi, Japanese, Mandarin) in Table 2 — models tend to be more literate when the information is asked in Western languages (see Table 6 in Appendix for the statistical test). A model coverage of languages in their pre-training data accounts for some of these trends; for instance, LLaMA 2 series underperforms in Hindi, and the Mixtral series weak results in Hindi, Japanese and Mandarin, as these languages are underrepresented in their respective training data. However, some trends are not explained. Mixtral, whose pre-training mainly focuses on English, German and Spanish, performs much better in Russian as compared to Hindi, Japanese or Mandarin which are also not focused on. Similarly, while LLaMA-3 claims to be more optimally trained in Hindi than Russian we note that Russian outperforms Hindi. Furthermore, GPT-4's highly varied performance in leaders and war-type questions across Western and non-Western languages. The inefficiency of these open-source models in handling non-Western languages raises significant concerns about inclusivity for a large and diverse user base (Blasi et al., 2021).

How transferable is cultural literacy of a model across languages? In the previous sections, we have analysed the cultural literacy of a model on individual languages. However, how much of this literacy is transferable across languages? For a model to be culturally literate across languages in a particular domain, it needs to consistently generate correct answers despite language variation. To evaluate the transference of cultural literacy of a model across languages, we first extract the number of factual entities a model can cover in a language, which we refer to as its ***coverage***. We define cov-

Model \ Lang	EN	DE	ES	HI	RU	JA	ZH	AVG	AVG$_W$	AVG$_{NW}$
Monuments										
Bloom-7B1	57.30	45.64	51.51	21.00	13.08	17.88	44.04	35.78	41.88	27.64
LLaMA-2-7B-chat	82.03	75.98	76.25	3.83	36.30	37.72	34.43	49.50	67.64	25.31
Mistral-7B-Instruct	88.26	80.07	81.41	29.27	65.30	26.25	33.54	57.73	78.76	29.68
Meta-LLaMA-3-8B-Instruct	59.16	88.52	87.90	53.91	57.38	55.43	63.08	66.48	73.24	57.46
LLaMA-2-13B-chat	85.41	77.76	80.16	15.39	54.80	37.19	31.49	54.60	74.53	28.02
Aya	64.15	62.81	59.52	35.68	30.34	34.34	32.83	45.67	54.20	34.29
GPT-4	**96.09**	**95.73**	**96.44**	**87.90**	**92.02**	**92.17**	**80.78**	**89.88**	**92.08**	**86.94**
Mixtral-8x7B	93.33	85.50	88.52	20.46	**82.12**	37.19	37.19	63.47	87.37	31.60
Leaders										
Bloom-7B1	22.25	26.83	20.08	21.67	24.00	26.58	26.00	23.92	23.29	24.76
LLaMA-2-7B-chat	56.42	46.58	27.42	8.50	23.00	28.33	28.00	31.18	38.35	21.62
Mistral-7B-Instruct	32.83	34.92	33.67	22.92	32.83	19.08	28.42	29.24	33.56	23.47
Meta-LLaMA-3-8B-Instruct	80.25	82.92	80.42	57.58	63.08	35.50	28.67	61.20	76.67	40.57
LLaMA-2-13B-chat	52.83	30.42	30.25	18.25	26.42	25.42	27.50	30.15	34.98	23.70
Aya	23.33	22.42	21.08	22.08	23.42	18.75	19.92	21.57	22.56	20.25
GPT-4	62.33	62.33	63.33	**47.33**	55.67	**48.00**	**48.33**	55.33	60.92	**47.87**
Mixtral-8x7B	**89.17**	**85.25**	**85.92**	9.33	**79.33**	32.75	28.08	**58.55**	**84.92**	23.39
Wars										
Bloom-7B1	32.93	25.11	29.04	33.96	28.12	28.79	39.24	31.02	28.80	33.98
LLaMA-2-7B-chat	56.27	46.07	42.25	24.93	34.49	39.27	35.94	39.89	44.77	33.38
Mistral-7B-Instruct	48.19	45.11	46.46	25.78	44.51	30.74	35.87	39.52	46.07	30.78
Meta-LLaMA-3-8B-Instruct	71.74	64.09	58.64	51.17	53.54	50.74	49.43	57.05	62.00	50.78
LLaMA-2-13B-chat	47.59	40.37	34.63	20.75	37.96	37.25	35.38	36.28	40.14	31.13
Aya	30.59	27.05	27.58	26.17	26.24	26.81	24.68	27.02	27.87	25.88
GPT-4	74.08	69.41	68.98	**57.79**	67.85	**60.34**	**59.07**	**65.36**	70.08	**59.06**
Mixtral-8x7B	**76.77**	**73.65**	**70.64**	26.27	66.29	49.04	47.52	58.60	**71.84**	40.94
National Parks										
Bloom-7B1	48.04	40.06	46.30	25.95	15.49	20.08	25.11	31.58	37.47	23.72
LLaMA-2-7B-chat	74.63	64.80	66.12	11.42	33.14	30.66	27.85	44.09	59.67	23.31
Mistral-7B-Instruct	77.33	72.94	67.39	19.50	51.11	13.90	25.69	46.84	67.19	19.70
Meta-LLaMA-3-8B-Instruct	64.22	79.60	80.02	51.32	56.77	39.11	40.96	58.86	70.15	43.80
LLaMA-2-13B-chat	75.53	68.97	66.70	10.15	42.49	31.92	26.27	46.01	63.42	22.79
Aya	50.32	45.30	44.71	36.21	33.30	32.66	24.42	38.13	43.41	31.09
GPT-4	**97.25**	**94.93**	**94.29**	**83.09**	**89.64**	**74.63**	**60.04**	**84.84**	**94.03**	**72.58**
Mixtral-8x7B	91.01	82.98	84.25	19.93	73.26	31.08	24.84	58.19	82.88	25.27

Table 2: Accuracy of LLMs across languages, averaged over countries. Column AVG_W and AVG_{NW} show the average performance of the models on Western (English, German, Spanish, and Russian) and Non-Western (Hindi, Japanese, and Mandarin) languages respectively. GPT-4 shows the best cultural literacy across all languages for almost all domains, with some exceptions. Mixtral-8x7B also has competitive results but fails to have high accuracy for non-Western languages (see Table 4 for ISO codes.).

erage, C_l^d, of a model in language l and domain d as the set of all entities whose questions it correctly answers a majority of the time, i.e., answers correctly in three out of four questions associated with that entity. Further, similar to Qi et al. (2023), we define the total coverage, TC^d, of a model over languages l as the area covered by the intersection of the coverage sets in domain d over their union on all languages as follows:

$$TC^d = \frac{|\bigcap_{l \in L} C_l^d|}{|\bigcup_{l \in L} C_l^d|} \qquad (1)$$

TC^d shows us the overall transference ability of the model. A model can have low accuracy but high TC^d, indicating better at transferring its knowledge. We evaluate TC^d in following scenarios:

1. **Total Coverage (All),** $TC^d(All)$**.** Here, we see the cultural transferability between all the languages XNationQA covers. This metric tells us the multilingual alignment ability of models on the dataset.

2. **Total Coverage (Pre-training data),** $TC^d(Pre - Train)$. Here, we see the cultural transferability between all the languages covered in model's training dataset. As discussed earlier some language are under-represented, particularly in LLaMA, Mistral AI models in the models pre-training language distribution. Therefore, to fairer judgment of these models, we also compute TC only on languages represented in training distribution of model.

3. **Total Coverage (West),** $TC^d(W)$**.** As highlighted by Table 2, models generally perform better in Western languages (English, German,

6

Model	TC^d (All)	TC^d (Pretrain)	TC^d (W)	TC^d (Eng-NW)
Monuments				
Bloomz-7b1	0.00	0.00	6.08	0.89
LLaMA-2-7b-chat	0.38	6.51	26.46	0.82
Mistral-7B-Instruct	6.34	78.16	57.20	7.34
Meta-Llama-3-8b-Instruct	14.18	30.26	34.59	18.95
LLaMA-2-13b-chat	2.27	8.71	44.87	2.40
Aya	6.01	6.01	21.90	10.70
GPT-4	62.59	62.59	81.88	72.56
Mixtral-8x7B	2.16	83.33	72.43	2.59
Leaders				
Bloomz-7b1	3.75	3.75	9.84	9.70
LLaMA-2-7b-chat	0.48	3.40	6.95	1.05
Mistral-7B-Instruct	2.29	54.62	41.54	4.03
Meta-Llama-3-8b-Instruct	10.25	44.80	48.75	12.17
LLaMA-2-13b-chat	6.67	14.12	24.71	8.28
Aya	16.95	16.95	30.69	22.45
GPT-4	28.46	28.46	50.83	33.48
Mixtral-8x7B	2.46	86.59	73.05	2.57
Wars				
Bloomz-7b1	5.21	5.21	11.39	13.50
LLaMA-2-7b-chat	3.47	9.31	16.67	8.53
Mistral-7B-Instruct	3.68	42.25	35.40	5.29
Meta-Llama-3-8b-Instruct	21.22	37.09	38.81	27.81
LLaMA-2-13b-chat	2.75	12.05	18.16	5.07
Aya	5.28	5.28	10.79	10.77
GPT-4	47.23	47.23	70.76	52.36
Mixtral-8x7B	9.45	71.26	59.33	11.30
National Parks				
Bloomz-7b1	0.29	0.29	7.72	0.65
LLaMA-2-7b-chat	0.48	5.60	21.54	0.52
Mistral-7B-Instruct	1.41	67.08	43.45	1.78
Meta-Llama-3-8b-Instruct	7.94	27.74	31.54	11.45
LLaMA-2-13b-chat	0.49	6.55	34.25	0.53
Aya	5.17	5.17	25.17	8.39
GPT-4	44.35	44.35	86.42	47.76
Mixtral-8x7B	1.08	77.88	60.35	1.11

Table 3: TC scores across XNationQA domains. Open-source models generally show low transferability.

Spanish, Russian). Therefore, we quantify their transferability in Western languages.

4. **Total Coverage (English Non-Western),** $TC^d(Eng - NW)$**.** As English is the dominant language for all models, we quantify how its transference is with respect to non-Western languages.

Table 3 shows the TC scores for all models in different domains. Except for GPT-4 and LLaMA-3, the $TC(All)$ score of all the models is near 5% — this shows the inability of current multilingual LLMs to be consistent with cultural knowledge across languages. Only about 5% of cultural literacy is truly transferable across languages. Interestingly, Aya appears to be better at transference compared to other models of similar size despite its overall accuracy being outstanding, as shown earlier. Also, as expected, the $TC(Pre - Train)$ score of models increases (on average, 10.4 ×) as compared to $TC(All)$.

For Western languages, TC is considerably higher than $TC(All)$. Mixtral and Mistral shows nearly a 30× and 18× jump in transference if we just consider the Western languages, this can

be attributed to its pretraining data as we saw a similar jump when comparing $TC(All)$ and $TC(Pre - Train)$ for these models. Interestingly, other models show stronger transference for $TC(W)$ than $TC(Pre - Train)$ indicating a biases towards Western languages, as also discussed in earlier section. This rise in $TC(W)$ can be attributed to the general bias these models exhibit towards Western languages (Naous et al., 2023). Also, $TC(Eng - NW)$ score for all models models is still low, highlighting the lack of alignment between English and non-Western languages.

We also analyse the transference of cultural literacy between languages. The pairwise TC (on all the domains) between languages pairs can be seen in Appendix A.5. Western language pairs have higher pairwise TC amongst themselves than other, going as high as 80 in some cases.

What disparities exist in a model's cultural literacy across nations? Figure 2 show the varying performance of the models on XNationQA. All models seem to have varying cultural literacy across languages and nations. Appendix Fig. 5 shows that Mixtral-8X7B performs most consistently across all nations for questions related to the birth year of leaders. However, its performance varies significantly across different languages, with non-Western languages lagging behind. This preference for Western languages over non-Western languages is evident in all models. LLaMA-3-8B-Instruct demonstrates a degree of consistent cultural literacy, particularly in English, German, and Spanish languages, though not on facts about Mexicans. LLaMA-2 7B and 13B variants show extreme cases where English outperforms all non-English languages; however, its performance varies widely for different nations. Surprisingly, GPT-4 shows varying performance across all languages, even English, and it is nowhere near as consistent as Mixtral-8XB. Additionally, no correlation is observed between a nation's most spoken language and the language in which the models show the best performance for questions about entities of that nation.

For questions related to the location of UNESCO sites (c.f. Appendix Fig. 6), we observe that, except for Aya and Bloomz, models generally show reasonable cultural literacy in Western languages. However, the LLaMA-2 family of models has low performance in facts related to Japan, Mexico, and China, even in the dominant language of these nations. For non-Western languages, models show

varying levels of performance. GPT-4 shows uniform cultural literacy across languages, even in Japanese and Hindi. However, other models exhibit non-uniform performance in these languages.

For questions on both the location of national parks (Appendix Fig. 8) and war years (Appendix Fig. 7), models exhibit highly variable literacy between nations and languages. GPT-4's performance in English, German, and Spanish for questions about national parks is the only notable exception.

To summarise our analysis, models generally perform well in Western languages but show varying levels of cultural literacy across different nations. With few exceptions, models generally lack uniform cultural literacy across languages. We also note that the model's best performance for facts about a nation is not always observed in the most popular language of that nation (Yin et al., 2022). These observations are especially true for the seven open-source models we experiment with, showcasing a wide gap in cultural literacy between open-source and closed-source models.

How literate are models about different nations? We find that GPT-4 consistently exhibits strong cultural literacy for these nations, particularly in the US, China, and India, where its accuracy is nearly 80%. LLaMA-3-8B-Instruct and Mixtral-8X7B also demonstrate decent performance, with accuracy reaching above 60% for the same countries. In contrast, the LLaMA-2 series shows moderate performance, in the 40–60% range. Aya and Bloomz show the lowest performance overall, with accuracy below 40% in most countries, indicating their limitations in handling cultural literacy. (Refer to Appendix Figure 13.)

Interestingly, the Western vs. Non-Western divide we consistently saw in other sections is not visible in model literacy. The models generally perform better in China, India, and Japan than in Germany, Russia, Mexico, and Spain, even though the models' performance in the native languages of these nations is lower (c.f. Table 2). This decoupling between linguistic competence and cultural knowledge underscores an important finding that a model may be culturally literate about a nation even if it underperforms in that nation's primary language.

6 Conclusion

In this paper, we demonstrated the inconsistency in cultural literacy of various LLMs in multiple
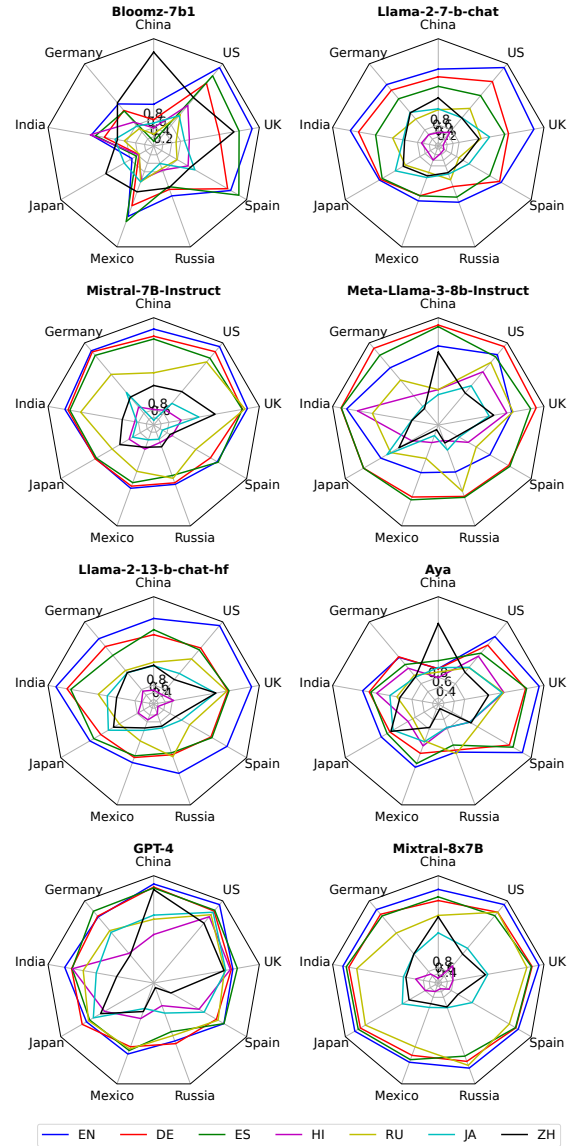
Figure 2: Accuracy of models in different languages for questions in XNationQA (see Table 4 for ISO code). We note the performance of all models vary across different language and nations.

languages. Our findings indicate that LLMs generally perform better in Western languages than non-Western ones. Models such as Aya and Bloom, designed for alignment between languages, perform poorly in recalling dates. Open-source models exhibit significant limitations in transferring knowledge across languages, with some models even having zero transferability scores in certain domains, highlighting the urgent need to address these disparities. Apart from a few instances of uniform performance, there is a noticeable variation in models' literacy levels regarding different nations, underscoring the biased coverage of topics from specific regions.

Limitations

This work attempts to present a more inclusive approach towards benchmarking the cultural literacy of models across nations and demonstrates the disparity in LLMs' performance with variation in language. However, due to computation limitations, we could not experiment with larger open-source LLMs. The cultural literacy of such large open-source multilingual LLMs is still unknown. Nevertheless, seven widely used multilingual open-source models have been benchmarked and the results indicate a wide variance of their literacy across languages.

While XNationQA consists of seven nations, our way of preferring to select highly spoken languages (Section A.3) resulted in nations from Africa and South America not being covered, This may be rectified by future works on the topic. Also, different nations have different numbers of national parks, UNESCO monuments, leaders or war instances — while XNationQA is parallel it contains different numbers of entities of each domain. This is an inherent limitation of benchmarking by focusing on facts about a nation.

We note that BLOOMZ and Aya, despite having more linguistically inclusive training data, are less culturally literate compared to other models. Analyzing the topic distribution and potential biases in these models' training datasets could give insights into this observation. However, since these datasets are not publicly available, we are unable to conduct such an analysis. Future work aimed at overcoming this limitation could be highly valuable for the development of multilingual models.

Ethical Statement

This work aims to evaluate the cultural literacy of large language models (LLMs) across multiple languages and nations by probing their knowledge of historically and culturally significant facts. We take care to ensure that our dataset is balanced, factual, and respectful of diverse cultures, avoiding stereotypes or biased representations. All cultural content is sourced from publicly available, reputable references such as encyclopedias, official historical records, and recognized heritage listings.

We acknowledge that cultural knowledge is complex and dynamic, and our work does not attempt to capture the full richness of any culture or community. Instead, it focuses on well-documented factual knowledge as a proxy for measuring models' cross-cultural understanding. We encourage future research to incorporate a wider range of cultural perspectives and to evaluate the social impact of deploying LLMs in multicultural contexts.

References

Julien Abadji, Pedro Ortiz Suarez, Laurent Romary, and Benoît Sagot. 2022. Towards a Cleaner Document-Oriented Multilingual Crawled Corpus. *arXiv e-prints*, arXiv:2201.06642.

Josh Achiam, Steven Adler, Sandhini Agarwal, Lama Ahmad, Ilge Akkaya, Florencia Leoni Aleman, Diogo Almeida, Janko Altenschmidt, Sam Altman, Shyamal Anadkat, et al. 2023. Gpt-4 technical report. *arXiv preprint arXiv:2303.08774*.

Kabir Ahuja, Rishav Hada, Millicent Ochieng, Prachi Jain, Harshita Diddee, Krithika Ramesh, Samuel Maina, Tanuja Ganu, Sameer Segal, Maxamed Axmed, Kalika Bali, and Sunayana Sitaram. 2023. Mega: Multilingual evaluation of generative ai. *ArXiv*, abs/2303.12528.

Shane Arora, Marzena Karpinska, Hung-Ting Chen, Ipsita Bhattacharjee, Mohit Iyyer, and Eunsol Choi. 2024. Calmqa: Exploring culturally specific long-form question answering across 23 languages. *arXiv preprint arXiv:2406.17761*.

Mikel Artetxe, Sebastian Ruder, and Dani Yogatama. 2019. On the cross-lingual transferability of monolingual representations. In *Annual Meeting of the Association for Computational Linguistics*.

Damián E. Blasi, Antonios Anastasopoulos, and Graham Neubig. 2021. Systematic inequalities in language technology performance across the world's languages. *ArXiv*, abs/2110.06733.

Terra Blevins, Hila Gonen, and Luke Zettlemoyer. 2022. Analyzing the mono- and cross-lingual pretraining dynamics of multilingual language models. *ArXiv*, abs/2205.11758.

Yong Cao, Li Zhou, Seolhwa Lee, Laura Cabello, Min Chen, and Daniel Hershcovich. 2023. Assessing cross-cultural alignment between chatgpt and human societies: An empirical study. *ArXiv*, abs/2303.17466.

Jonathan H Clark, Eunsol Choi, Michael Collins, Dan Garrette, Tom Kwiatkowski, Vitaly Nikolaev, and Jennimaria Palomaki. 2020. Tydi qa: A benchmark for information-seeking question answering in typologically di verse languages. *Transactions of the Association for Computational Linguistics*, 8:454–470.

Viet Dac Lai, Chien Van Nguyen, Nghia Trung Ngo, Thuat Nguyen, Franck Dernoncourt, Ryan A Rossi, and Thien Huu Nguyen. 2023. Okapi: Instruction-tuned large language models in multiple languages with reinforcement learning from human feedback. *arXiv e-prints*, pages arXiv–2307.

9

Ameet Deshpande, Partha Talukdar, and Karthik Narasimhan. 2022. When is BERT multilingual? isolating crucial ingredients for cross-lingual transfer. In *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 3610–3623, Seattle, United States. Association for Computational Linguistics.

Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. 2018. Bert: Pre-training of deep bidirectional transformers for language understanding. *arXiv preprint arXiv:1810.04805*.

Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aiesha Letman, Akhil Mathur, Alan Schelten, Amy Yang, Angela Fan, et al. 2024. The llama 3 herd of models. *arXiv preprint arXiv:2407.21783*.

Stefan Daniel Dumitrescu, Petru Rebeja, Beata Lorincz, Mihaela Gaman, Andrei Avram, Mihai Ilie, Andrei Pruteanu, Adriana Stan, Lorena Rosia, Cristina Iacobescu, Luciana Morogan, George Dima, Gabriel Marchidan, Traian Rebedea, Madalina Chitez, Dani Yogatama, Sebastian Ruder, Radu Tudor Ionescu, Razvan Pascanu, and Viorica Patraucean. 2021. Liro: Benchmark and leaderboard for romanian language tasks. In *Thirty-fifth Conference on Neural Information Processing Systems Datasets and Benchmarks Track (Round 1)*.

Fahim Faisal, Yinkai Wang, and Antonios Anastasopoulos. 2021. Dataset geography: Mapping language data to language users. *arXiv preprint arXiv:2112.03497*.

Yi Fung, Ruining Zhao, Jae Doo, Chenkai Sun, and Heng Ji. 2024. Massively multi-cultural knowledge acquisition & lm benchmarking. *arXiv preprint arXiv:2402.09369*.

Rishav Hada, Varun Gumma, Adrian de Wynter, Harshita Diddee, Mohamed Ahmed, Monojit Choudhury, Kalika Bali, and Sunayana Sitaram. 2023. Are large language model-based evaluators the solution to scaling up multilingual evaluation? In *Findings*.

William B. Held, Camille Harris, Michael Best, and Diyi Yang. 2023. A material lens on coloniality in nlp. *ArXiv*, abs/2311.08391.

Erik D Hirsch. 1983. Cultural literacy. *The American Scholar*, pages 159–169.

Albert Q Jiang, Alexandre Sablayrolles, Arthur Mensch, Chris Bamford, Devendra Singh Chaplot, Diego de las Casas, Florian Bressand, Gianna Lengyel, Guillaume Lample, Lucile Saulnier, et al. 2023. Mistral 7b. *arXiv preprint arXiv:2310.06825*.

Albert Q Jiang, Alexandre Sablayrolles, Antoine Roux, Arthur Mensch, Blanche Savary, Chris Bamford, Devendra Singh Chaplot, Diego de las Casas, Emma Bou Hanna, Florian Bressand, et al. 2024. Mixtral of experts. *arXiv preprint arXiv:2401.04088*.

Zhengbao Jiang, Antonios Anastasopoulos, Jun Araki, Haibo Ding, and Graham Neubig. 2020. X-FACTR: Multilingual factual knowledge retrieval from pretrained language models. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 5943–5959, Online. Association for Computational Linguistics.

Nora Kassner, Philipp Dufter, and Hinrich Schütze. 2021. Multilingual LAMA: Investigating knowledge in multilingual pretrained language models. In *Proceedings of the 16th Conference of the European Chapter of the Association for Computational Linguistics: Main Volume*, pages 3250–3258, Online. Association for Computational Linguistics.

Amr Keleg and Walid Magdy. 2023. Dlama: A framework for curating culturally diverse facts for probing the knowledge of pretrained language models. *arXiv preprint arXiv:2306.05076*.

Tom Kwiatkowski, Jennimaria Palomaki, Olivia Redfield, Michael Collins, Ankur Parikh, Chris Alberti, Danielle Epstein, Illia Polosukhin, Jacob Devlin, Kenton Lee, et al. 2019. Natural questions: a benchmark for question answering research. *Transactions of the Association for Computational Linguistics*, 7:453–466.

Shayne Longpre, Yi Lu, and Joachim Daiber. 2021. Mkqa: A linguistically diverse benchmark for multilingual open domain question answering. *Transactions of the Association for Computational Linguistics*, 9:1389–1406.

Mai Miyabe and Takashi Yoshino. 2015. Evaluation of the validity of back-translation as a method of assessing the accuracy of machine translation. *2015 International Conference on Culture and Computing (Culture Computing)*, pages 145–150.

Junho Myung, Nayeon Lee, Yi Zhou, Jiho Jin, Rifki Putri, Dimosthenis Antypas, Hsuvas Borkakoty, Eunsu Kim, Carla Perez-Almendros, Abinew Ali Ayele, et al. 2024. Blend: A benchmark for llms on everyday knowledge in diverse cultures and languages. *Advances in Neural Information Processing Systems*, 37:78104–78146.

Tarek Naous, Michael Joseph Ryan, and Wei Xu. 2023. Having beer after prayer? measuring cultural bias in large language models. *ArXiv*, abs/2305.14456.

Thuat Nguyen, Chien Van Nguyen, Viet Dac Lai, Hieu Man, Nghia Trung Ngo, Franck Dernoncourt, Ryan A Rossi, and Thien Huu Nguyen. 2023. Culturax: A cleaned, enormous, and multilingual dataset for large language models in 167 languages. *arXiv preprint arXiv:2309.09400*.

Jirui Qi, Raquel Fern'andez, and Arianna Bisazza. 2023. Cross-lingual consistency of factual knowledge in multilingual language models. In *Conference on Empirical Methods in Natural Language Processing*.

Swarnadeep Saha, Omer Levy, Asli Celikyilmaz, Mohit Bansal, Jason Weston, and Xian Li. 2023. Branch-solve-merge improves large language model evaluation and generation. *ArXiv*, abs/2310.15123.

Pola Schwöbel, Jacek Golebiowski, Michele Donini, Cedric Archambeau, and Danish Pruthi. 2023. Geographical erasure in language generation. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pages 12310–12324, Singapore. Association for Computational Linguistics.

Weiyan Shi, Ryan Li, Yutong Zhang, Caleb Ziems, Raya Horesh, Rogério Abreu de Paula, Diyi Yang, et al. 2024. Culturebank: An online community-driven knowledge base towards culturally aware language technologies. *arXiv preprint arXiv:2404.15238*.

Margaret S Steffensen, Chitra Joag-Dev, and Richard C Anderson. 1979. A cross-cultural perspective on reading comprehension. *Reading research quarterly*, pages 10–29.

Hugo Touvron, Louis Martin, Kevin Stone, Peter Albert, Amjad Almahairi, Yasmine Babaei, Nikolay Bashlykov, Soumya Batra, Prajjwal Bhargava, Shruti Bhosale, et al. 2023. Llama 2: Open foundation and fine-tuned chat models. *arXiv preprint arXiv:2307.09288*.

Ahmet Üstün, Viraat Aryabumi, Zheng-Xin Yong, Wei-Yin Ko, Daniel D'souza, Gbemileke Onilude, Neel Bhandari, Shivalika Singh, Hui-Lee Ooi, Amr Kayid, et al. 2024. Aya model: An instruction finetuned open-access multilingual language model. *arXiv preprint arXiv:2402.07827*.

Linting Xue, Noah Constant, Adam Roberts, Mihir Kale, Rami Al-Rfou, Aditya Siddhant, Aditya Barua, and Colin Raffel. 2020. mt5: A massively multilingual pre-trained text-to-text transformer. *arXiv preprint arXiv:2010.11934*.

Linting Xue, Noah Constant, Adam Roberts, Mihir Kale, Rami Al-Rfou, Aditya Siddhant, Aditya Barua, and Colin Raffel. 2021. mT5: A massively multilingual pre-trained text-to-text transformer. In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 483–498, Online. Association for Computational Linguistics.

Da Yin, Hritik Bansal, Masoud Monajatipoor, Liunian Harold Li, and Kai-Wei Chang. 2022. Geomlama: Geo-diverse commonsense probing on multilingual pre-trained language models. In *EMNLP*.

Zheng Xin Yong, Hailey Schoelkopf, Niklas Muennighoff, Alham Fikri Aji, David Ifeoluwa Adelani, Khalid Almubarak, M Saiful Bari, Lintang Sutawika, Jungo Kasai, Ahmed Baruwa, Genta Winata, Stella Biderman, Edward Raff, Dragomir Radev, and Vassilina Nikoulina. 2023. BLOOM+1: Adding language support to BLOOM for zero-shot prompting. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 11682–11703, Toronto, Canada. Association for Computational Linguistics.

Kaitlyn Zhou, Kawin Ethayarajh, and Dan Jurafsky. 2022. Richer countries and richer representations. In *Findings of the Association for Computational Linguistics: ACL 2022*, pages 2074–2085, Dublin, Ireland. Association for Computational Linguistics.

A Appendix

A.1 Prompt Templates

Table 5 shows a few question templates used in our dataset creation. We translate the question template along with the entity-answer tuples and options to create XNationQA.

A.2 ISO Code

Table 4 contains the ISO codes of the seven languages used in our dataset.

Language	ISO 639-1 code	Family
English	EN	IE: Germanic
German	DE	IE: Germanic
Hindi	HI	IE: Indo-Iranian
Mandarin	ZH	Sino-Tibetan
Russian	RU	IE: Balto-Slavic
Spanish	ES	IE: Italic
Japanese	ZH	Japonic

Table 4: List of languages and their ISO codes used in our experiments.

A.3 Selecting languages and entity distribution

We started our study by selecting the three most widely spoken languages in the world (Mandarin, Spanish and English). One nation from each of these languages (China, Spain, US) was selected in the beginning, we then expanded our selection to include a combination of diverse (Hindi, Russian) and similar languages (German and Japanese), with one nation where they are widely spoken (India, Russia, Germany), to these languages. This resulted in a set of seven nations and languages, we included Mexico and UK to expand on the cultures we cover. Figure 3 shows the distribution of entities over the nine countries.

906
907
908
909
910
911
912
913
914
915
916
917
918
919
920
921
922
923
924
925
926
927
928
929
930
931
932
933
934
935
936
937

A.4 Training dataset of base models of the models used in our study.

Aya: Aya is an instruction-tuned mT5 model(Xue et al., 2020) that supports 101 languages, including all seven languages used in our study.

Bloomz models: Bloomz supports 46 languages, including programming languages. It covers all seven languages used in our study.

Mistral models: The Mistral family of models primarily focuses on European languages, covering English, French, Italian, German, and Spanish. The exact distribution of its training dataset is unknown, so its coverage of other languages is unknown.

LLaMA-2 models: LLaMA-2's pre-training dataset predominantly consists of English, with 89.7% of the data in English. The remaining 8.38% comprises unknown languages (mainly programming code), and about 2% consists of non-English languages. Except for Hindi, all other languages in our dataset seem to be represented.

Meta-LLaMA-3-8B series: LLaMA-3 builds upon the multilingual capabilities of LLaMA-2, with 8% of its training dataset dedicated to multiple languages. Although LLaMA-3 has been trained on a wide array of languages, it is specifically optimized and safely tuned for eight languages: English, German, French, Italian, Portuguese, Hindi, Spanish, and Thai.

GPT-4: The exact details of its training dataset are not known publicaly.

Figure 3: Distribution of the entities across nations.

938
939
940
941
942
943
944
945
946
947
948
949
950
951
952
953
954
955
956
957
958
959
960
961
962
963
964
965
966
967
968
969
970
971
972
973
974
975
976
977
978
979
980
981
982
983
984
985

A.5 Pairwise TC of models

Figure 4 represents the pairwise TC of Bloomz-7b1, LLaMA-2-7-b-chat, Mistral-7b-instruct, LLaMA-2-13-b-chat, and Aya. We see that the model generally sees a better TC amount in Western languages compared to non-Western languages, highlighting that the models struggle in non-Western languages. Mistral-7b-instruct and LLaMA-2-13-b-chat show relatively better performance among the models evaluated, but none as strong as GPT-4, Mixtral or Meta-LLaMA-3 that have been discussed in the main text.

LLaMA-2 series shows very poor TC for all pairs consisting of Hindi, indicating the lack of consistency and understanding in Hindi. Aya while showing low performance in Figure 2 shows overall reasonable TC across languages.

A.6 Evaluating disparity across national park and war questions

Figure 7 shows the performance of models on war-type questions. We note that models have different literacy on these facts with no model performing similarly across all countries or languages. Dates about wars Russia has taken part in seem to be difficult for the model to recall, except for Mixtral and Meta-LLaMA-3.

Fig. 11 shows the performance of models on the location of national park-type questions. Here also we note that models tend to have different literacy on these facts with no model performing similarly. Except for GPT-4 for Western languages, there is no homogeneous performance across countries.

A.7 Model-wise analysis of results

Here we dive deeper into analysing each model for a specific use case and along the country-language axes as shown in figure 9,10,11 and 12. The analysis is as follows:

Bloomz-7b1: We note that for war and leader domain questions, the model shows modest to poor in most cases, even worse than random in some instances. Despite being trained on a diverse multilingual corpus, Bloomz struggles with `XNationQA`compared to LLaMA or Mistral models of the same size. For questions related to monuments and national parks, the model performs better in English, German, Spanish, and Chinese but shows low performance in other languages.

12

Domain	Question Template	Entity-Answer example
Year of War	Question: As per American history, when did {War} happen? {Option} Answer: Question: What year did {War} start according to American history? {Option} Answer: Question: Historical accounts from American indicate {War} occur in? {Option} Answer: Question: Can you tell me according to American history the date when {War} took place? {Option} Answer:	*(American Revolutionary War, 1775)*
Location of national park or monument	Question: Where can {Location} be found within China? {Option} Answer: Question: In which part of China is {Location} situated? {Option} Answer: Question: Where in China is {Location} located? {Option} Answer: Question: In which area of China can {Location} be found? {Option} Answer:	*(Wuyi Mountains National Park, Fujian)*
Birth year of Leader	'Question: When was German leader {NAME} born? {Option} Answer: Question: What is the birthyear of the German leader {NAME}? {Option} Answer: Question: On what year was the German leader {NAME} born? {Option} Answer: Question: What is the year of birth of the German leader {NAME}? {Option} Answer:	*(Olaf Scholz,1958)*

Table 5: Question Template in English along with an example entity-answer. We create the option using three negative samples and the true answer to fill the {Option}.

Model	M	L	W	N
Bloom-7B1	**1.9e-47**	0.38	**0.002**	**1.23e-59**
LLaMA-2-7B-chat	**2.82e-160**	**4.16e-39**	**1.20e-38**	**2.83e-195**
Mistral-7B-Instruct	**4.96e-113**	**1.324e-5**	**3.06e-243**	**3.77e-189**
Meta-LLaMA-3-8B	**1.145e-31**	**3.90e-86**	**1.82e-32**	**1.02e-81**
LLaMA-2-13B-chat	**1.83e-32**	**7.35e-15**	**9.82e-12**	**4.67e-193**
Aya	**3.73e-45**	0.45	0.02	**1.06e-21**
Mixtral-8x7B	**1.09e-132**	**2.67e-156**	**1.47e-105**	**3.99e-247**
GPT-4	**1.25e-13**	**4.23e-9**	**2.45e-12**	**3.57e-51**

Table 6: p-values for significance testing of the hypothesis of Modeling performing better in Western languages. M (Monuments), L (Leader), W (War) and N (National Park). Significance is decided by taking $\alpha = 0.05$ and has been marked bold

LLaMA-2-7-b-chat: We consistently see that in all domains the model performs badly in the Hindi language (With accuracy lower than 20%). It seems to be most literate in Western-European languages which consistently outperform other languages for all countries and domains. Its coverage in Russian is still moderate compared to Western European languages. While it does show a decent level of literacy about a nation in the nation's native language They are not the nest language to prompt the model in.

Mistral-7b-instruct: It shows modest performance on leader domain questions, but for other domains, we note it to be culturally literate. However, it is biased towards Western languages and shows only modest performance in other languages, especially Hindi.

Meta-LLaMA-3-8B-Instruct: It is the only model below the size of 10 billion to be literate to some degree in all the languages covered by us, even Hindi. It still seems to be more proficient in Western languages than others.

LLaMA-2-13-b-chat: Similar to the 7 billion variant of LLaMA this model struggles to perform well in Hindi and has considerable disparity between Western-European languages and other languages.

Aya: Like Bloomz, Aya has been trained on a wide range of diverse multilingual corpus, and sim-
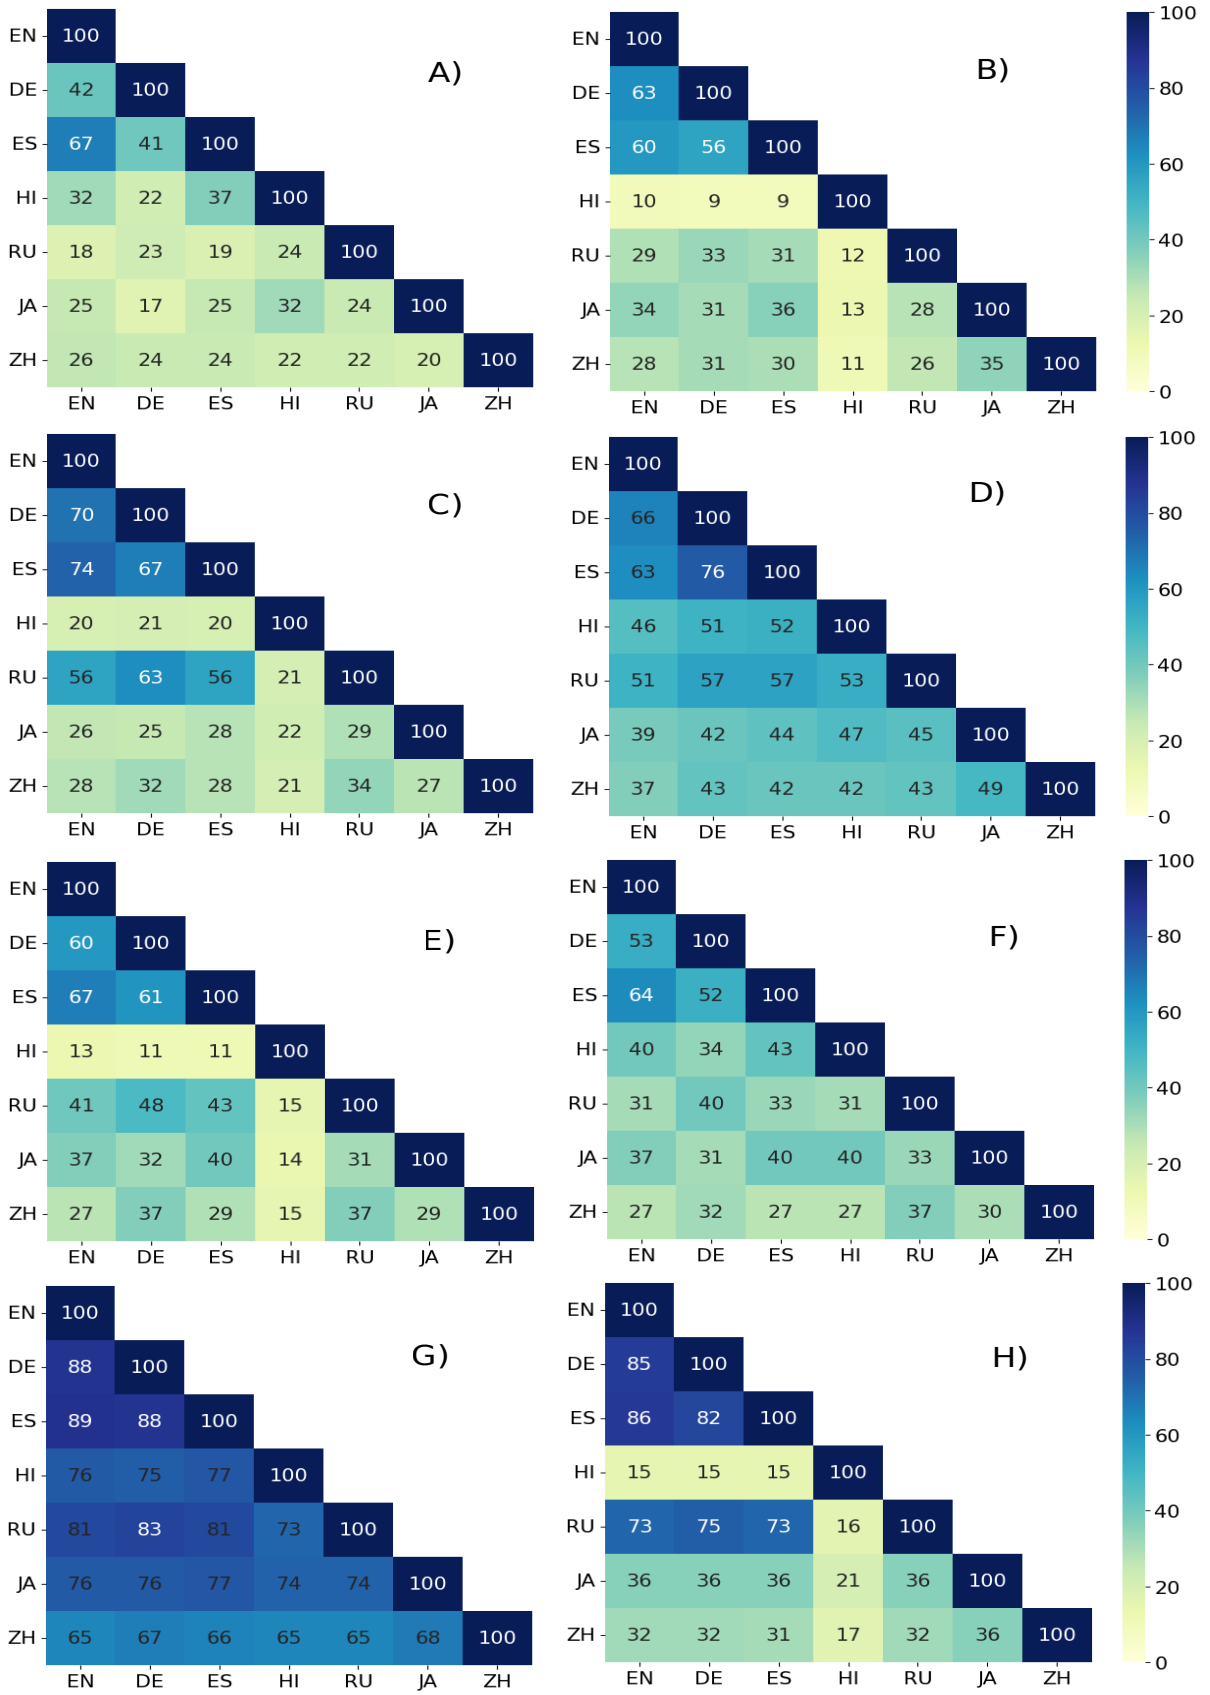
13

Figure 4: Heat map pairwise TC of all language pairs – (A) Bloomz-7b1, (B)LLaMA-2-7-b-chat, (C)Mistral-7b-instruct, (D)Meta-LLaMA-3-8B-Instruct, (E) LLaMA-2-13-b-chat, (F) 13-billion Aya, (G) GPT-4 and (H) Mixtral-8X7B.
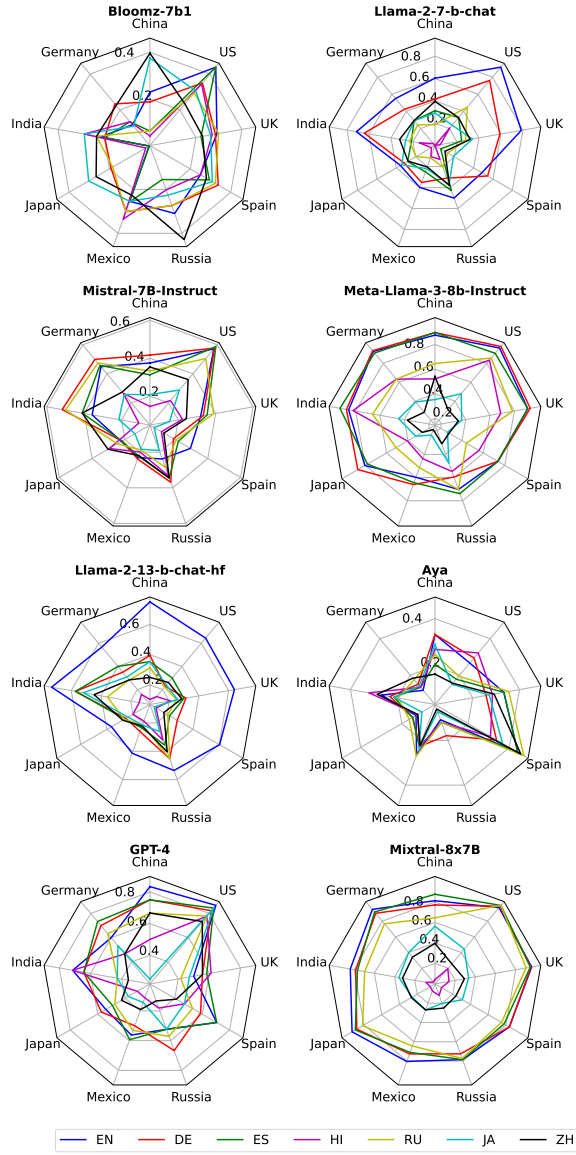
Figure 5: Accuracy of models in different languages for questions on the birth year of leaders of different nations (see Table 4 for ISO code).
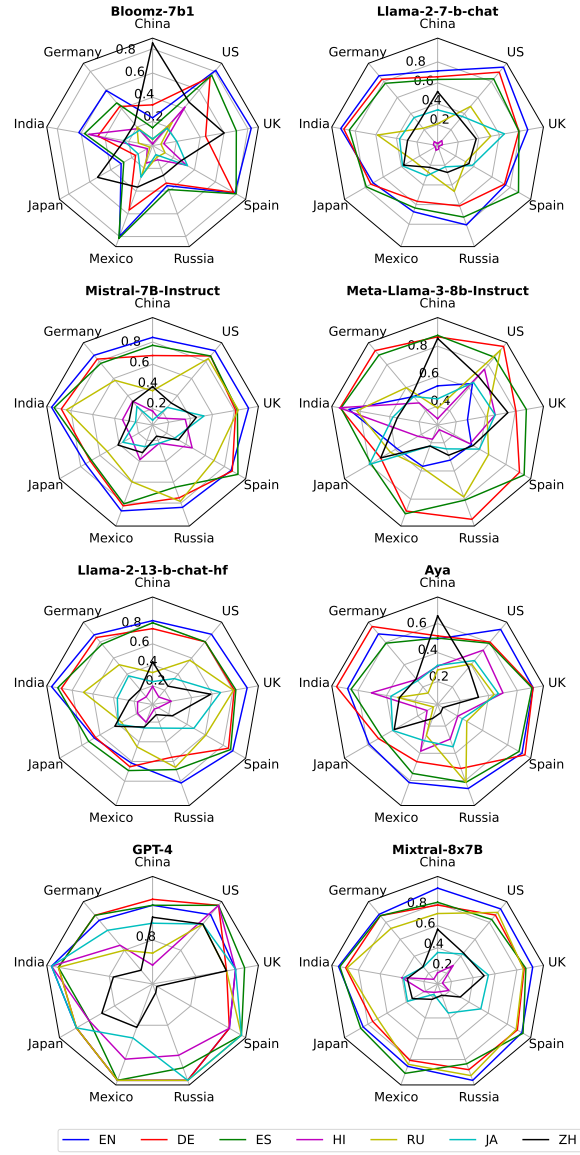

Figure 6: Accuracy of models in different languages for questions on the location of UNESCO sites of different nations (see Table 4 for ISO code).

ilar to Bloomz it struggles to perform well in war and leader domain questions. But its performance in other domains is better, but not comparable to LLaMA-2-13-b-chat which is of a similar size.

GPT-4: It is the most culturally literate model in multilingual setups. With considerably higher performance than all other models. There is an interesting observation though it struggles with questions about Germany if asked in Russian in monuments and national park domain.

B Human Evaluation of dataset.

To evaluate the dataset, we recited university students aged 20 to 25 who were either native speakers of the language or had cleared language proficiency texts. We recruited a total of six language experts, one for each language. They were provided with the following Annotation Guidelines and Scoring Criteria.

Annotation Guidelines: Each translation is evaluated based on the following criteria:

- **Grammar** – Proper use of verbs, tense, agreement, and punctuation.

- **Fluency** – The sentence should sound natural and idiomatic.

- **Cohesion** – The sentence should be logically structured and clear.
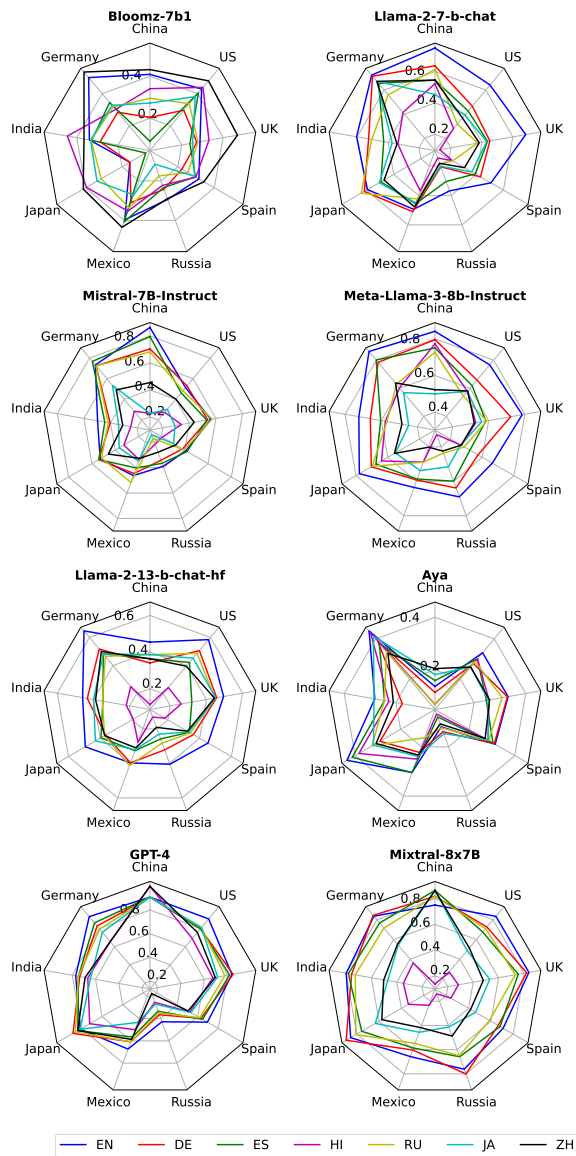
15

Figure 7: Accuracy of models in different languages (indicated in different colours) for questions on the date of wars in different nations.


Figure 8: Accuracy of models in different languages (indicated in different colours) for questions on the location of national parks of different nations.

Scoring Criteria: Each translation is rated on a scale from 1 to 5:

- **5 (Excellent)** – No grammatical errors. The sentence follows correct conjugations and agreements. It is natural, idiomatic, and fluent, resembling native speech.

- **4 (Good)** – Minor grammatical mistakes (e.g., incorrect prepositions, small conjugation errors), but they do not affect understanding. The sentence is still readable and sounds natural.

- **3 (Acceptable)** – Noticeable grammar mistakes, such as verb tense inconsistencies or incorrect word order. The sentence is understandable but slightly unnatural.

- **2 (Poor)** – Multiple grammatical mistakes significantly affecting readability. The sentence sounds unnatural or contains awkward phrasing.

- **1 (Unacceptable)** – The sentence is not grammatically correct and is difficult or impossible to understand. Severe errors break fluency, making the translation unusable.

16

C Evaluating the variation of performance of the model across languages.

To study the disparity in a model's performance between languages on facts about the same country, we analyze the standard deviation of the accuracy of the models in different languages. A lower standard deviation indicates better homogeneous performance across languages. However, we observe high standard variance (c.f. Table 7), indicating that multilingual models struggle to transfer their cultural literacy knowledge through the languages. In many instances, the standard deviation exceeds twenty, indicating a severe disparity in performance transference across languages. We will elaborate this further.

Country / Model	US	China	India	Uk	Japan	Germany	Russia	Mexico	Spain
Monuments									
Bloomz-7b1	51.93 ± 23.55	24.05 ± 26.89	39.80 ± 18.07	42.08 ± 28.19	21.07 ± 15.99	32.30 ± 15.57	23.38 ± 11.96	46.29 ± 24.76	48.47 ± 26.86
LLaMA-2-7B-chat	56.85 ± 32.05	44.17 ± 24.04	57.06 ± 32.36	57.70 ± 26.98	47.14 ± 26.27	48.25 ± 31.01	44.81 ± 26.47	38.87 ± 21.98	48.98 ± 28.92
Mistral-7B-Instruct	59.97 ± 31.97	47.52 ± 27.06	63.69 ± 32.38	67.63 ± 21.04	53.04 ± 17.28	56.98 ± 24.03	51.95 ± 27.87	58.10 ± 24.20	63.61 ± 25.67
Meta-LLaMA-3-8B-Instruct	78.42 ± 12.33	61.08 ± 22.62	80.36 ± 15.75	69.20 ± 12.21	66.61 ± 12.08	57.30 ± 24.21	61.36 ± 21.45	59.34 ± 20.28	67.35 ± 16.40
LLaMA-2-13B-chat	54.02 ± 28.50	50.80 ± 25.36	61.65 ± 32.55	69.31 ± 22.56	48.93 ± 18.09	53.33 ± 29.10	45.29 ± 27.55	44.64 ± 19.47	57.82 ± 30.42
Aya	52.38 ± 11.52	43.73 ± 12.96	51.11 ± 16.86	55.13 ± 14.45	36.79 ± 18.11	42.22 ± 22.37	44.97 ± 19.10	39.01 ± 15.25	42.01 ± 26.08
GPT-4	95.83 ± 3.86	85.42 ± 8.05	95.92 ± 7.17	92.86 ± 2.19	92.14 ± 3.64	79.68 ± 21.43	93.51 ± 10.83	93.96 ± 7.65	93.20 ± 11.05
Mixtral-8x7B	67.86 ± 25.99	62.03 ± 25.21	70.24 ± 27.99	67.30 ± 27.96	58.75 ± 24.89	57.94 ± 31.36	62.50 ± 35.23	57.14 ± 33.89	67.69 ± 30.52
Leader									
Bloomz-7b1	34.37 ± 6.99	19.35 ± 13.76	23.47 ± 3.50	26.59 ± 2.77	13.93 ± 11.87	14.64 ± 5.54	26.70 ± 8.15	27.57 ± 3.50	29.20 ± 3.19
LLaMA-2-7B-chat	46.11 ± 25.30	31.55 ± 17.59	36.22 ± 20.71	37.20 ± 22.42	25.88 ± 10.29	30.24 ± 14.99	33.16 ± 11.90	23.49 ± 10.17	21.64 ± 19.71
Mistral-7B-Instruct	44.44 ± 14.23	30.65 ± 8.89	35.97 ± 13.88	28.77 ± 7.77	23.42 ± 5.12	38.33 ± 10.22	30.10 ± 6.09	21.15 ± 1.99	17.65 ± 6.85
Meta-LLaMA-3-8B-Instruct	76.03 ± 24.05	67.56 ± 20.10	70.92 ± 20.47	70.44 ± 23.54	56.62 ± 22.94	67.26 ± 24.52	58.16 ± 14.15	46.82 ± 17.78	53.57 ± 19.57
LLaMA-2-13B-chat	29.05 ± 15.14	36.01 ± 18.25	47.70 ± 17.38	32.14 ± 13.35	26.64 ± 4.65	33.57 ± 11.67	39.80 ± 8.53	25.00 ± 7.01	24.37 ± 15.70
Aya	22.30 ± 6.50	26.49 ± 5.76	24.74 ± 4.21	30.36 ± 3.40	13.41 ± 1.89	14.76 ± 2.21	11.05 ± 3.58	23.94 ± 1.82	40.34 ± 5.93
GPT-4	80.95 ± 4.42	63.10 ± 18.29	56.12 ± 13.99	52.78 ± 6.30	45.67 ± 8.52	59.05 ± 9.87	50.34 ± 10.45	46.21 ± 10.36	55.46 ± 10.83
Mixtral-8x7B	69.37 ± 34.96	58.63 ± 26.87	55.10 ± 26.99	64.58 ± 34.36	58.67 ± 34.88	61.43 ± 33.31	52.04 ± 28.85	50.00 ± 26.65	54.62 ± 30.31
Wars									
Bloomz-7b1	39.58 ± 5.78	29.22 ± 10.96	34.69 ± 4.55	30.47 ± 7.25	26.33 ± 12.41	37.32 ± 9.33	22.45 ± 5.84	35.87 ± 5.36	28.91 ± 3.66
LLaMA-2-7B-chat	41.31 ± 10.99	58.12 ± 8.83	44.50 ± 9.50	41.22 ± 15.10	51.66 ± 10.83	63.05 ± 10.14	22.85 ± 7.31	45.50 ± 4.39	35.18 ± 9.75
Mistral-7B-Instruct	42.42 ± 9.40	56.49 ± 24.37	37.89 ± 7.93	46.68 ± 10.50	46.84 ± 8.23	60.84 ± 16.54	27.10 ± 9.14	40.53 ± 6.70	34.99 ± 7.03
Meta-LLaMA-3-8B-Instruct	61.78 ± 6.96	70.13 ± 12.47	58.13 ± 8.80	61.30 ± 9.38	65.78 ± 7.56	69.83 ± 10.63	51.93 ± 12.07	53.34 ± 7.37	52.13 ± 6.86
LLaMA-2-13B-chat	44.16 ± 9.90	34.42 ± 9.98	37.11 ± 7.29	41.03 ± 7.96	37.46 ± 8.56	47.78 ± 10.86	25.96 ± 8.24	34.86 ± 4.88	33.25 ± 7.40
Aya	28.81 ± 2.35	13.31 ± 4.92	23.81 ± 3.58	28.43 ± 3.45	35.38 ± 5.23	38.79 ± 4.07	11.17 ± 2.72	24.53 ± 4.54	29.26 ± 1.53
GPT-4	79.54 ± 6.02	93.51 ± 4.11	71.43 ± 4.41	76.53 ± 6.61	82.39 ± 4.79	78.33 ± 9.53	31.07 ± 7.77	56.83 ± 7.30	59.33 ± 6.72
Mixtral-8x7B	60.23 ± 17.63	73.05 ± 24.55	62.64 ± 16.80	63.45 ± 19.24	69.44 ± 18.05	66.01 ± 16.65	55.90 ± 21.42	50.08 ± 13.36	53.80 ± 15.92
National Parks									
Bloomz-7b1	43.15 ± 22.78	16.43 ± 19.95	29.96 ± 10.43	54.76 ± 35.01	16.70 ± 9.41	24.76 ± 10.25	27.56 ± 9.69	29.69 ± 13.11	48.66 ± 32.61
LLaMA-2-7B-chat	52.94 ± 27.82	38.57 ± 23.41	47.12 ± 25.45	53.81 ± 30.65	46.95 ± 21.53	36.19 ± 17.65	32.38 ± 14.22	31.81 ± 15.33	45.54 ± 27.04
Mistral-7B-Instruct	59.50 ± 25.86	44.29 ± 29.93	49.14 ± 29.51	79.52 ± 27.58	40.23 ± 20.14	47.62 ± 25.69	36.96 ± 17.38	32.14 ± 20.54	37.28 ± 27.26
Meta-LLaMA-3-8B-Instruct	66.01 ± 11.78	60.00 ± 22.99	62.09 ± 20.36	79.29 ± 8.16	63.03 ± 15.82	36.90 ± 21.67	51.63 ± 18.15	45.54 ± 19.77	53.79 ± 15.49
LLaMA-2-13B-chat	53.40 ± 28.04	48.57 ± 19.77	47.92 ± 27.22	70.48 ± 24.44	56.09 ± 20.98	36.90 ± 14.65	32.61 ± 15.85	34.71 ± 18.00	38.62 ± 23.32
Aya	42.86 ± 13.39	24.29 ± 25.13	40.44 ± 9.90	59.29 ± 15.68	39.29 ± 7.23	22.86 ± 7.90	30.36 ± 10.76	31.92 ± 8.06	32.14 ± 20.66
GPT-4	93.32 ± 4.98	85.71 ± 23.21	87.30 ± 15.09	98.10 ± 3.01	87.82 ± 5.08	66.67 ± 33.24	73.29 ± 18.62	76.56 ± 19.30	83.04 ± 17.90
Mixtral-8x7B	63.42 ± 24.37	52.86 ± 28.27	60.48 ± 32.37	72.38 ± 23.67	60.71 ± 28.82	49.05 ± 23.48	50.00 ± 29.42	49.00 ± 29.10	58.93 ± 34.04

Table 7: Accuracy and standard divination of accuracy on the seven languages covered by XNationQA on the countries covered. All models show high variation in accuracy over languages. Results with standard deviation higher than ten are marked with red color.
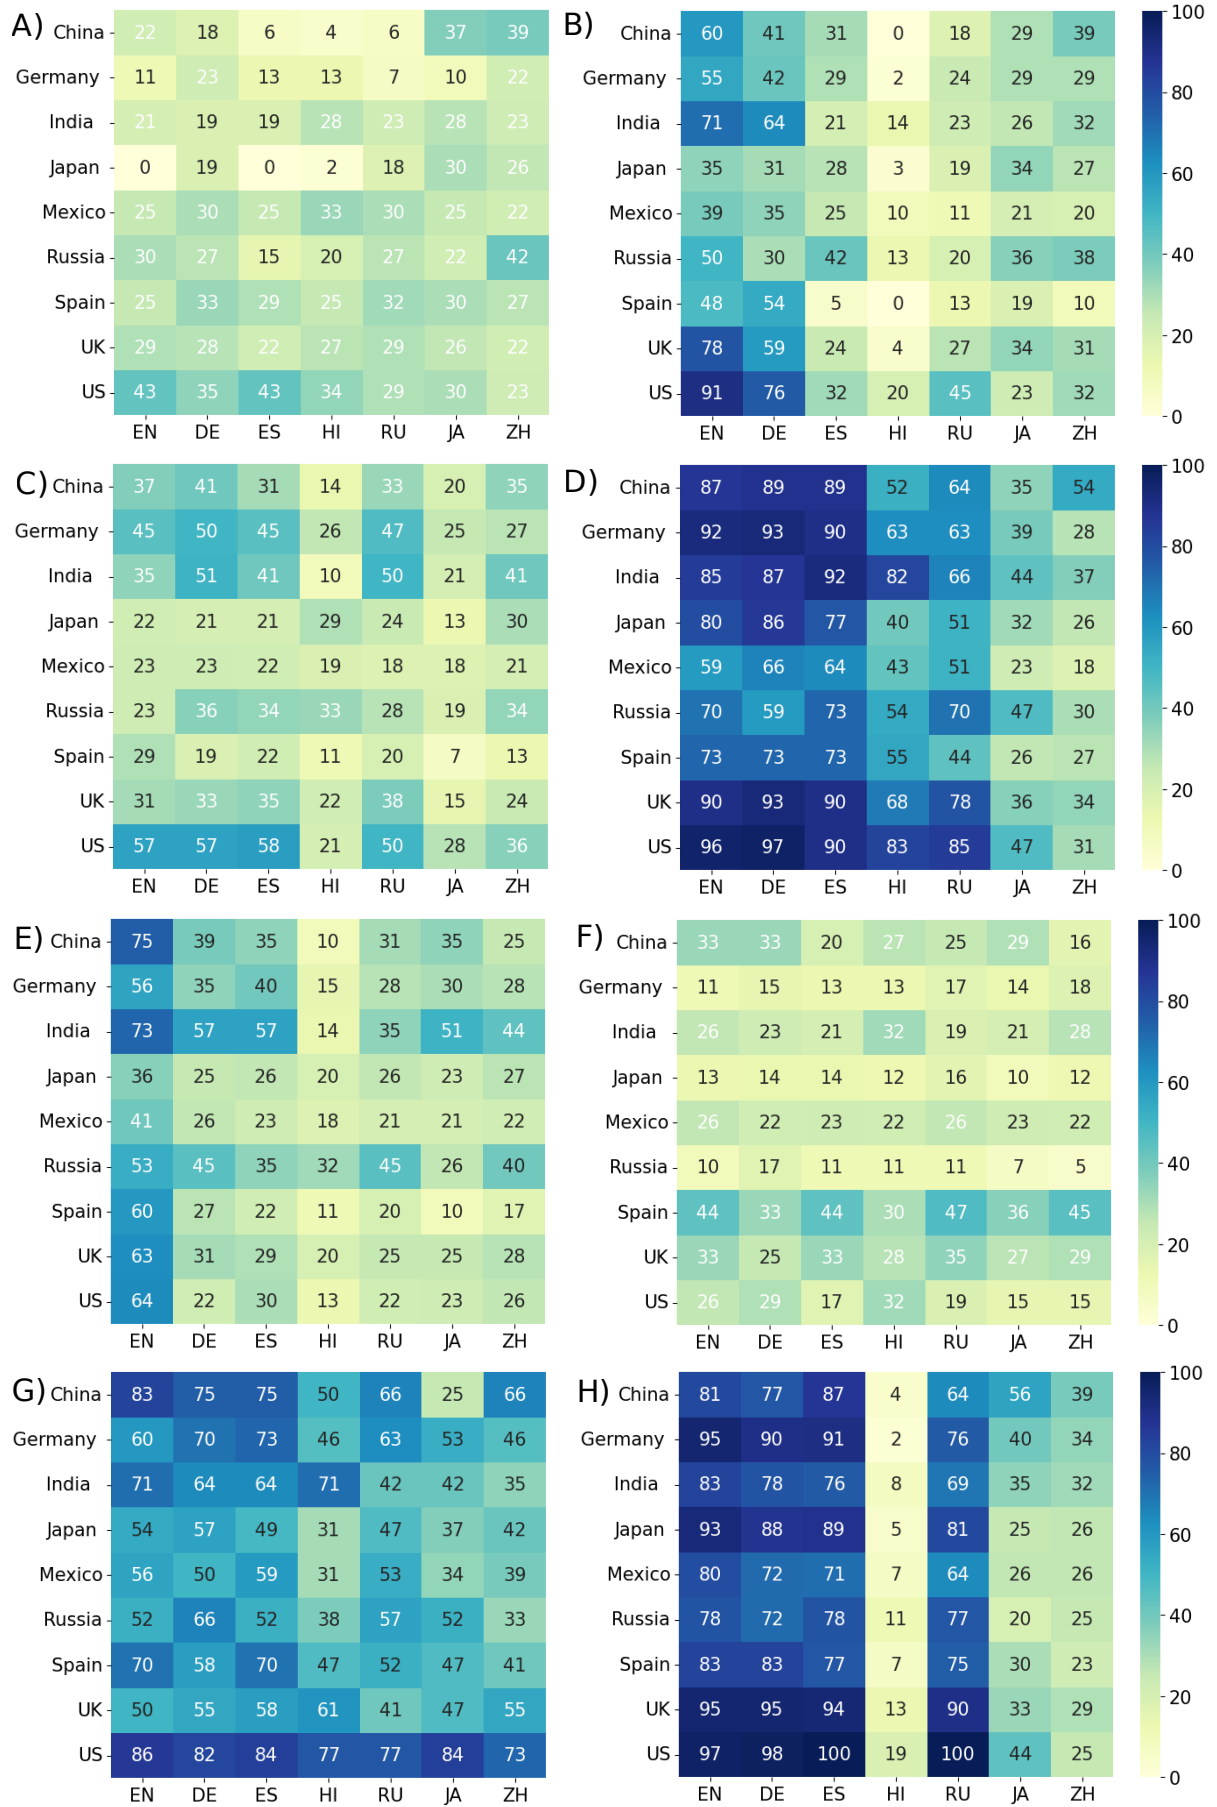
Figure 9: Heat map of accuracy for leader domain A) Bloomz-7b1, B)LLaMA-2-7-b-chat C)Mistral-7b-instruct D)Meta-LLaMA-3-8B-Instruct E) LLaMA-2-13-b-chat and F) 13-billion Aya G) GPT-4 H) Mixtral-8X7B.
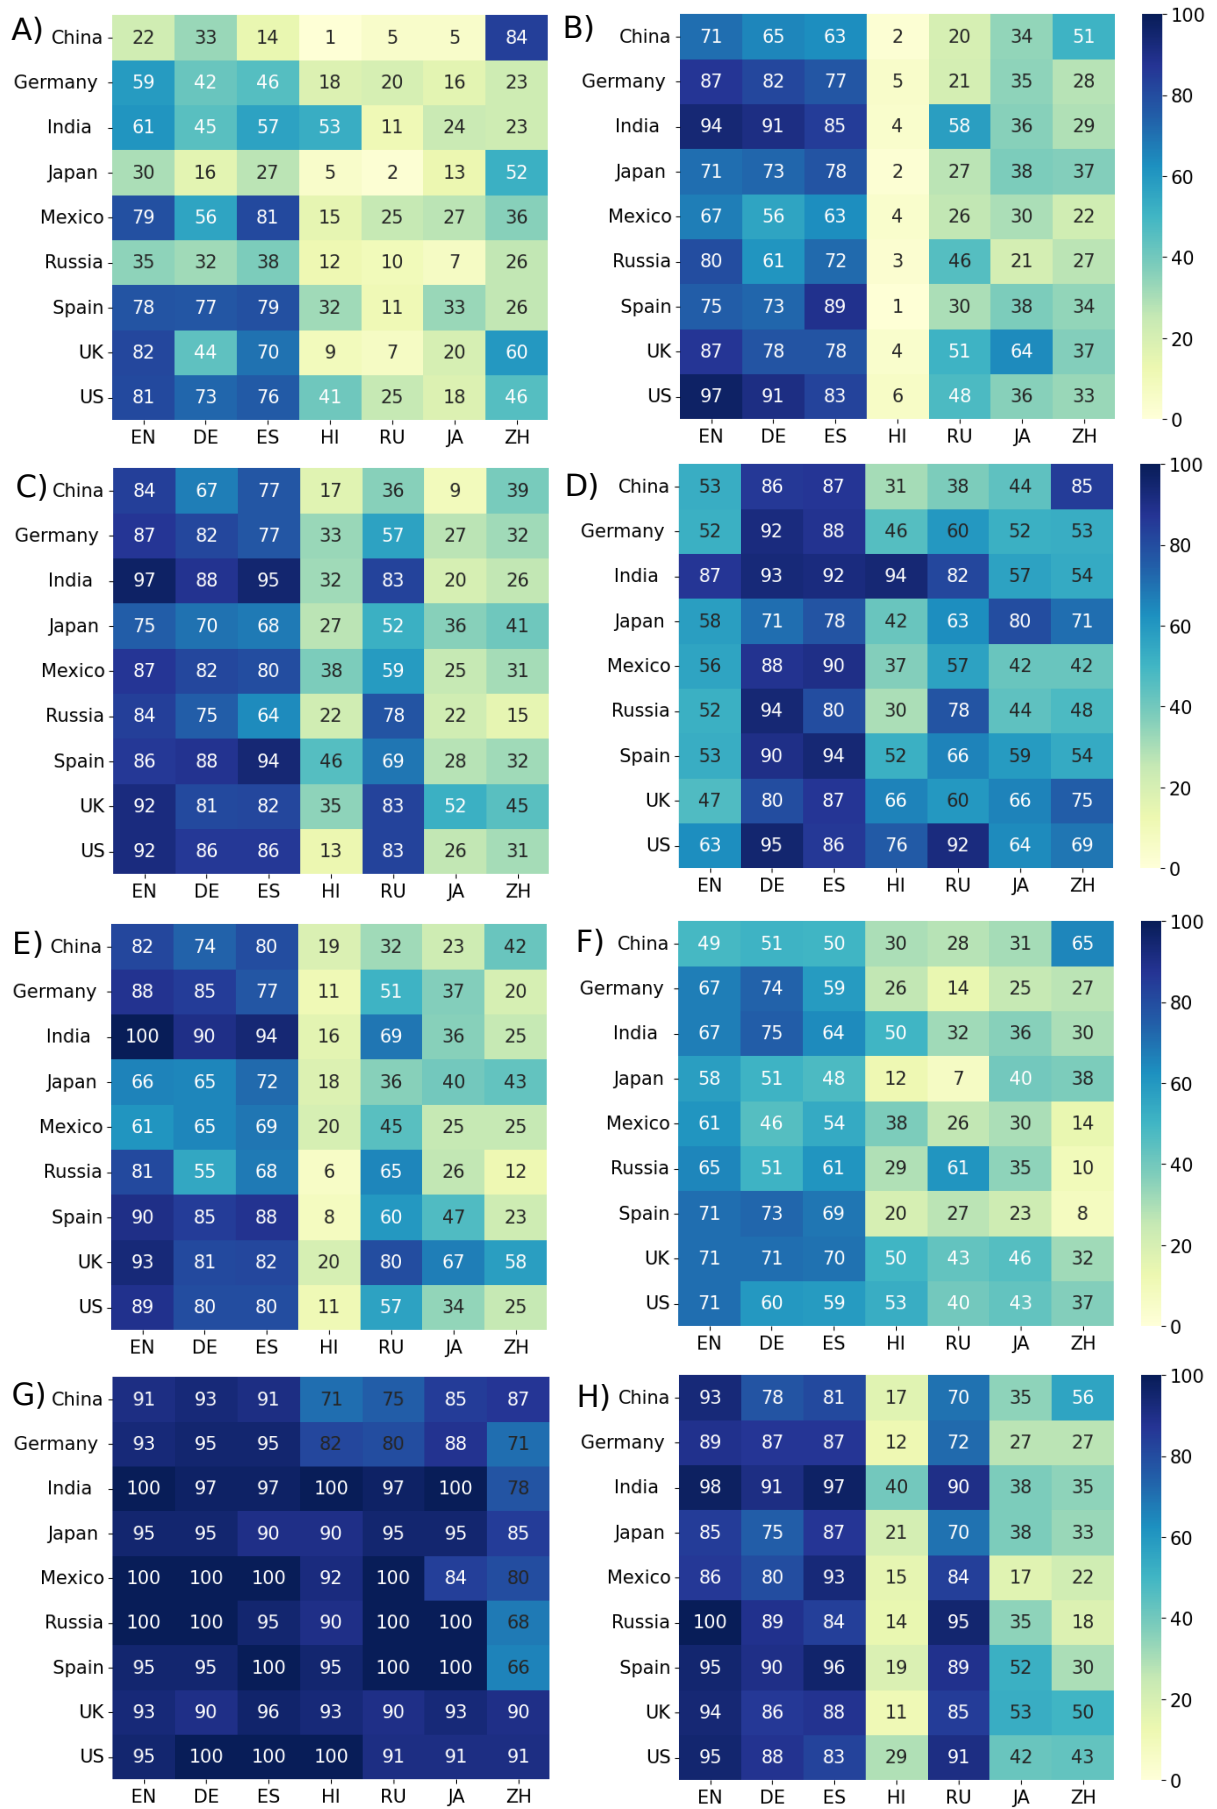
Figure 10: Heat map of accuracy for monuments domain: (A) Bloomz-7b1, (B)LLaMA-2-7-b-chat (C) Mistral-7b-instruct, (D) Meta-LLaMA-3-8B-Instruct, (E) LLaMA-2-13-b-chat, and (F) 13-billion Aya, (G) GPT-4 and (H) Mixtral-8X7B.
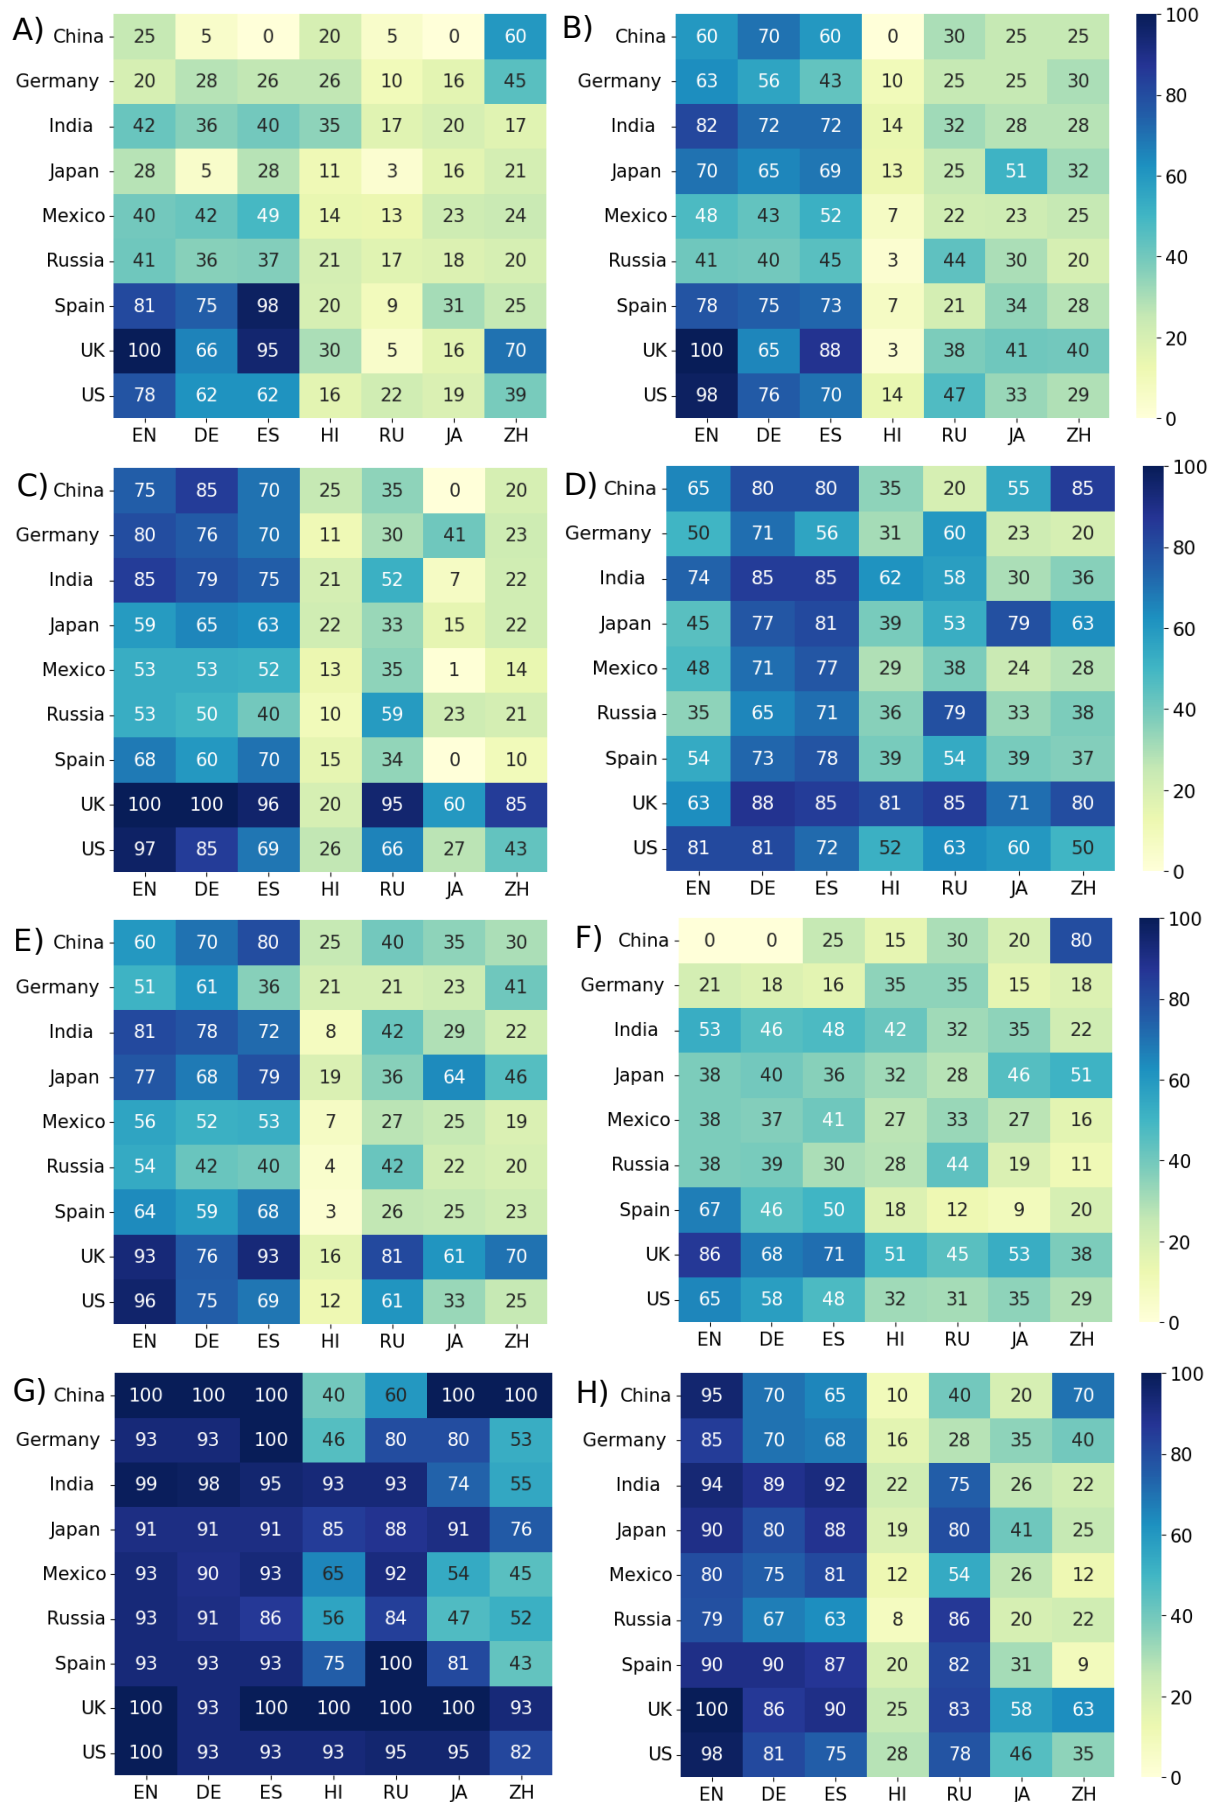
Figure 11: Heat map of accuracy for national park domain: A) Bloomz-7b1, B)LLaMA-2-7-b-chat C)Mistral-7b-instruct D)Meta-LLaMA-3-8B-Instruct E) LLaMA-2-13-b-chat and F) 13-billion Aya G) GPT-4 H) Mixtral-8X7B.

Figure 12: Heat map of accuracy for war domain: A) Bloomz-7b1, B)LLaMA-2-7-b-chat C)Mistral-7b-instruct D)Meta-LLaMA-3-8B-Instruct E) LLaMA-2-13-b-chat and F) 13-billion Aya G) GPT-4 H) Mixtral-8X7B.

Figure 13: Accuracy of models across nations averaged over language. We note that models like GPT have lesser performance in Mexico, Spanish, and Russian as compared to India and China even though the native languages of these nations (Hindi and Mandarin) are poorly performing (Table 2).